

z/OS 3.1 IBM Education Assistant

Solution Name: RSM Dedicated Real Memory Pools >4T for specific functions

Solution Element(s): BCP/RSM, BCP/zCX, BCP/SVA

July 2023



Agenda

- Trademarks
- Objectives
- Overview
- Usage & Invocation
- Interactions & Dependencies
- Upgrade & Coexistence Considerations
- Installation & Configuration
- Summary
- Appendix

Trademarks

- See url <http://www.ibm.com/legal/copytrade.shtml> for a list of trademarks.
- Additional Trademarks:
 - None

Objectives

- Facilitate the usage of >4T real memory
- Assign memory directly to applications instead of first come first serve
 - Avoid system disruptions due to memory spikes or irregular memory usage patterns

Overview

- Who is interested:
 - Installations that are concerned about applications with irregular or unpredictable memory usage
 - SVC Dump capture
 - Installations that want to preferentially assign memory to certain applications that exploit high virtual storage
 - zCX containers
 - Installations that want to exploit >4T of memory
- What (Solution)
 - Provide the ability for the system administrator to define an amount of real storage that can be used to be directly assigned to job steps
 - Provide the ability to assign Dedicated Memory in 2G units directly to applications and support its use transparently
 - Provide diagnostics aids to determine how effectively the Dedicated Memory is used
- Wow (Benefit / Value, Need Addressed)
 - Utilization of >4T for applications that do not exploit 2G frames
 - Mitigate disruptions from SVC DUMP
 - Better control over memory

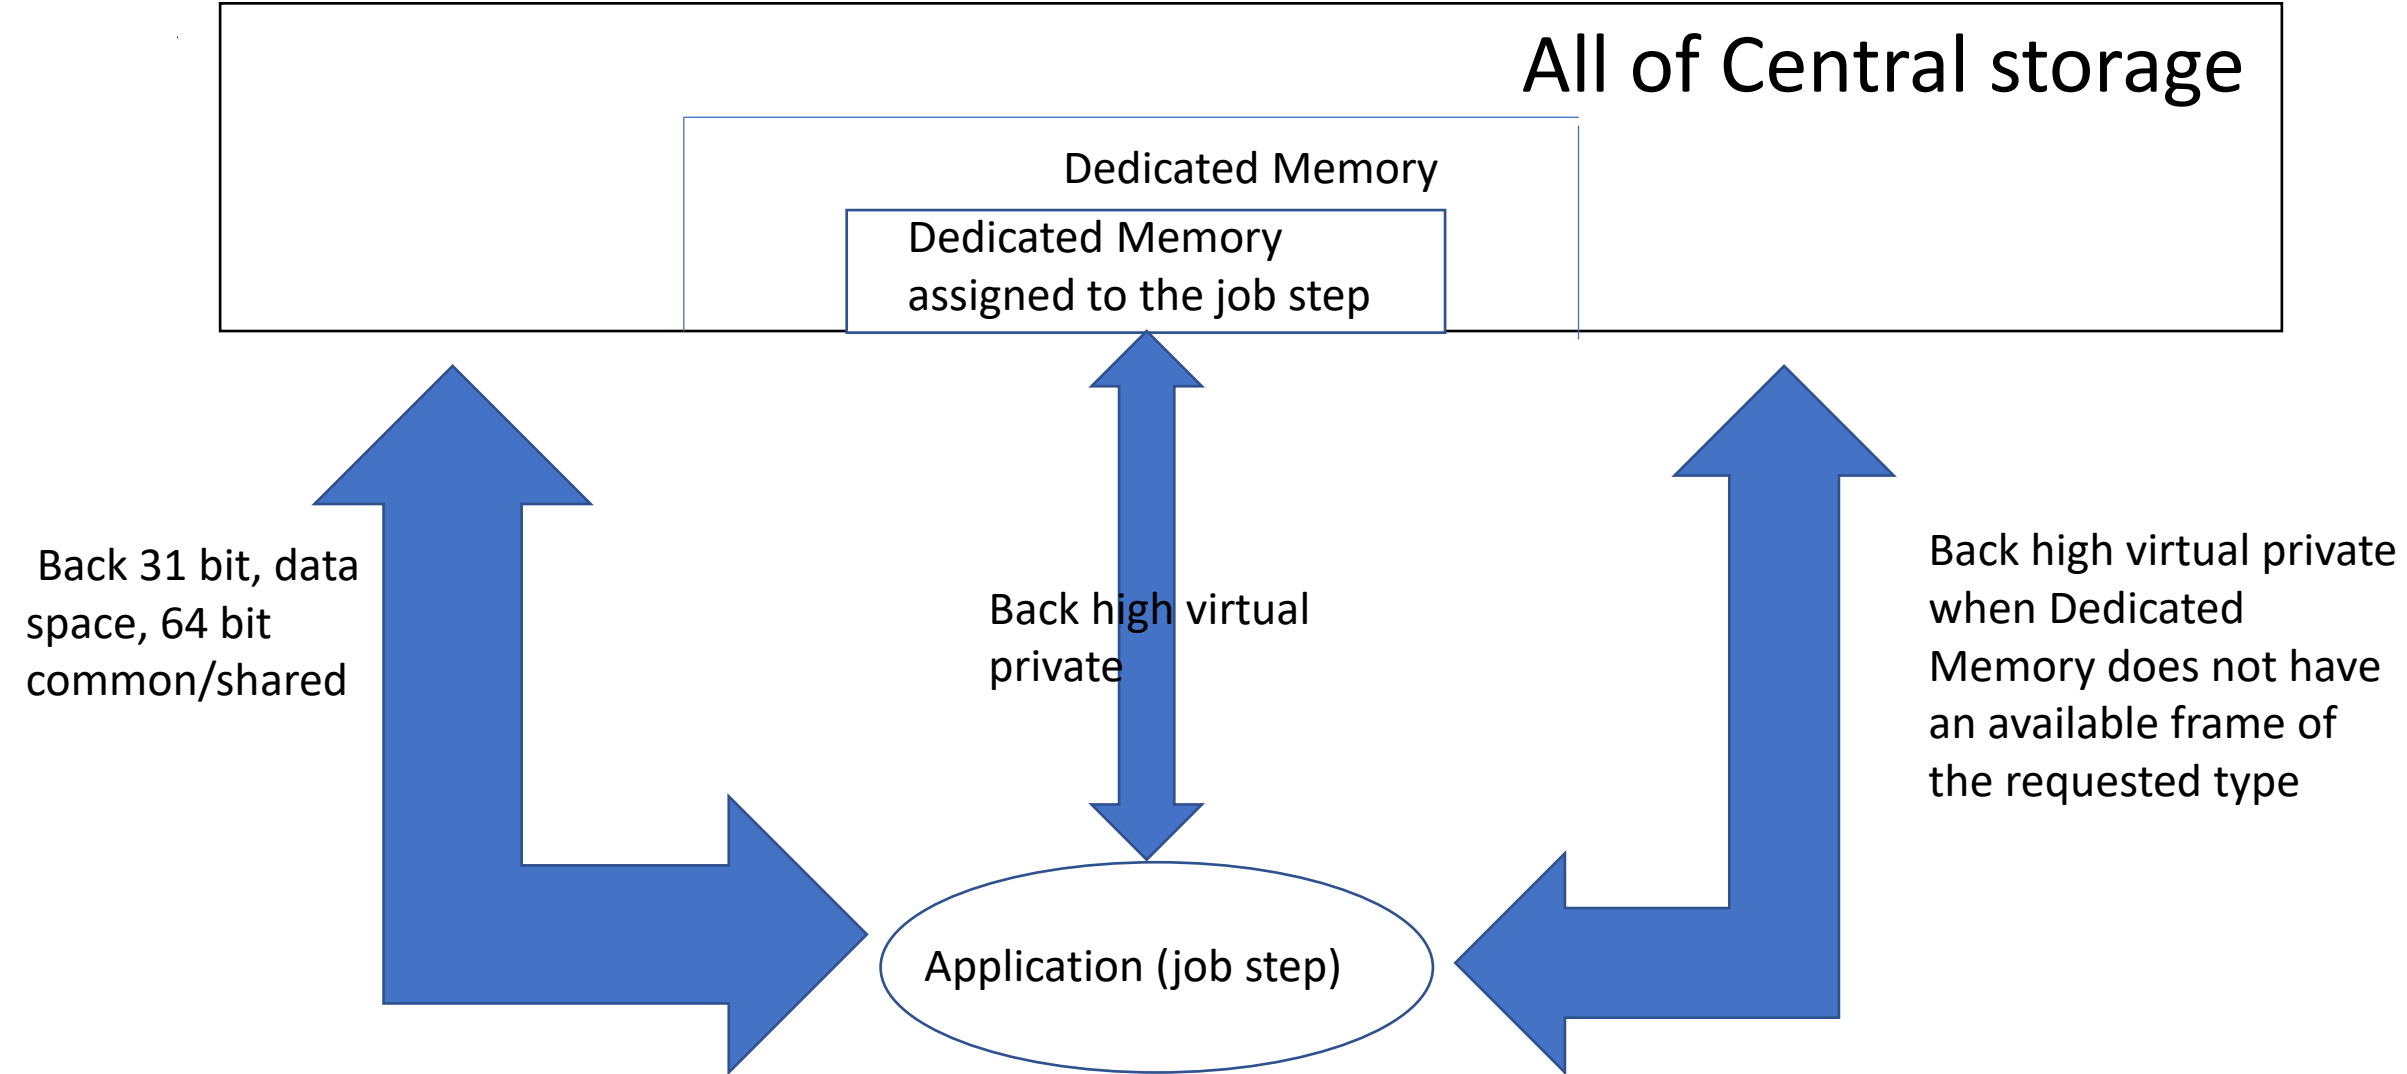
Dedicated Memory Usage

- Is used (transparently) back any private memory object as long as the memory object is freed at end of job step i.e. owned by the job step task or a descendent
 - Is used transparently to back page tables of such memory objects
- Is used to back any private memory object if the address space is a single step started task (also any DAT tables)
 - DUMPSRV
 - zCX
- New SMF30 and SMF71 fields for instrumentation (see appendix)
- Never paged/stolen – address space owns the memory until end of job step, regardless of whether it actually uses the memory
- Dedicated Memory assignment not incorporated into SYSEVENT STGTEST result
- Not counted against MEMPOOL

Dedicated Memory Usage

- Not used for 31 bit memory (data spaces, hiperspaces, 31 bit address space)
- Not used for common or shared high virtual
- Not used for paging I/O
- Not included in existing frame counts
- Cannot be used for OMVS

Dedicated Memory Usage



Interactions & Dependencies

- Software Dependencies
 - None
- Hardware Dependencies
 - None
- Exploiters
 - Exploitation is up to the installation, however, the following applications are recommended
 - zCX container instances
 - DUMPSRV

Upgrade & Coexistence Considerations

- To exploit this solution, all systems in the Plex must be at the new z/OS level: No
- Toleration/coexistence APARs/PTFs.
 - None

Installation and Configuration

- In order to define Dedicated Memory to the system, the new parmlib member IARPRMxx needs to be specified with the new DEDICATEDMEMORY parameter:
 - DEDICATEDMEMORY(xxG|xxT)
- New IEASYSxx parm RSM=
 - RSM=(xx,yy) – where xx and yy correspond to IARPRMxx and IARPRMyy
- The size of the reserved dedicated memory area cannot be dynamically changed after IPL.
- Rules for DEDICATEDMEMORY
 - Must have at least 16G of online NON-Dedicated Memory after the Dedicated Memory area is designated
 - Total system storage must be a multiple of 2G
 - Must be specified as a 2G multiple
 - System will round up or down to a storage increment multiple if necessary

Initialization & IPL Considerations

- How much Dedicated Memory to designate on an image?
 - DUMPSRV – DISPLAY DUMP,INFO can be used to obtain historical sizes of SVC dumps for the current ipl
 - zCX – Already requires knowing how much memory to assign
 - Should not exceed Memlimit
- Some Dedicated Memory is used to manage Dedicated Memory
 - For every 256G of assignable Dedicated Memory, the system needs 2G to manage it
 - Define more Dedicated Memory than you plan to assign

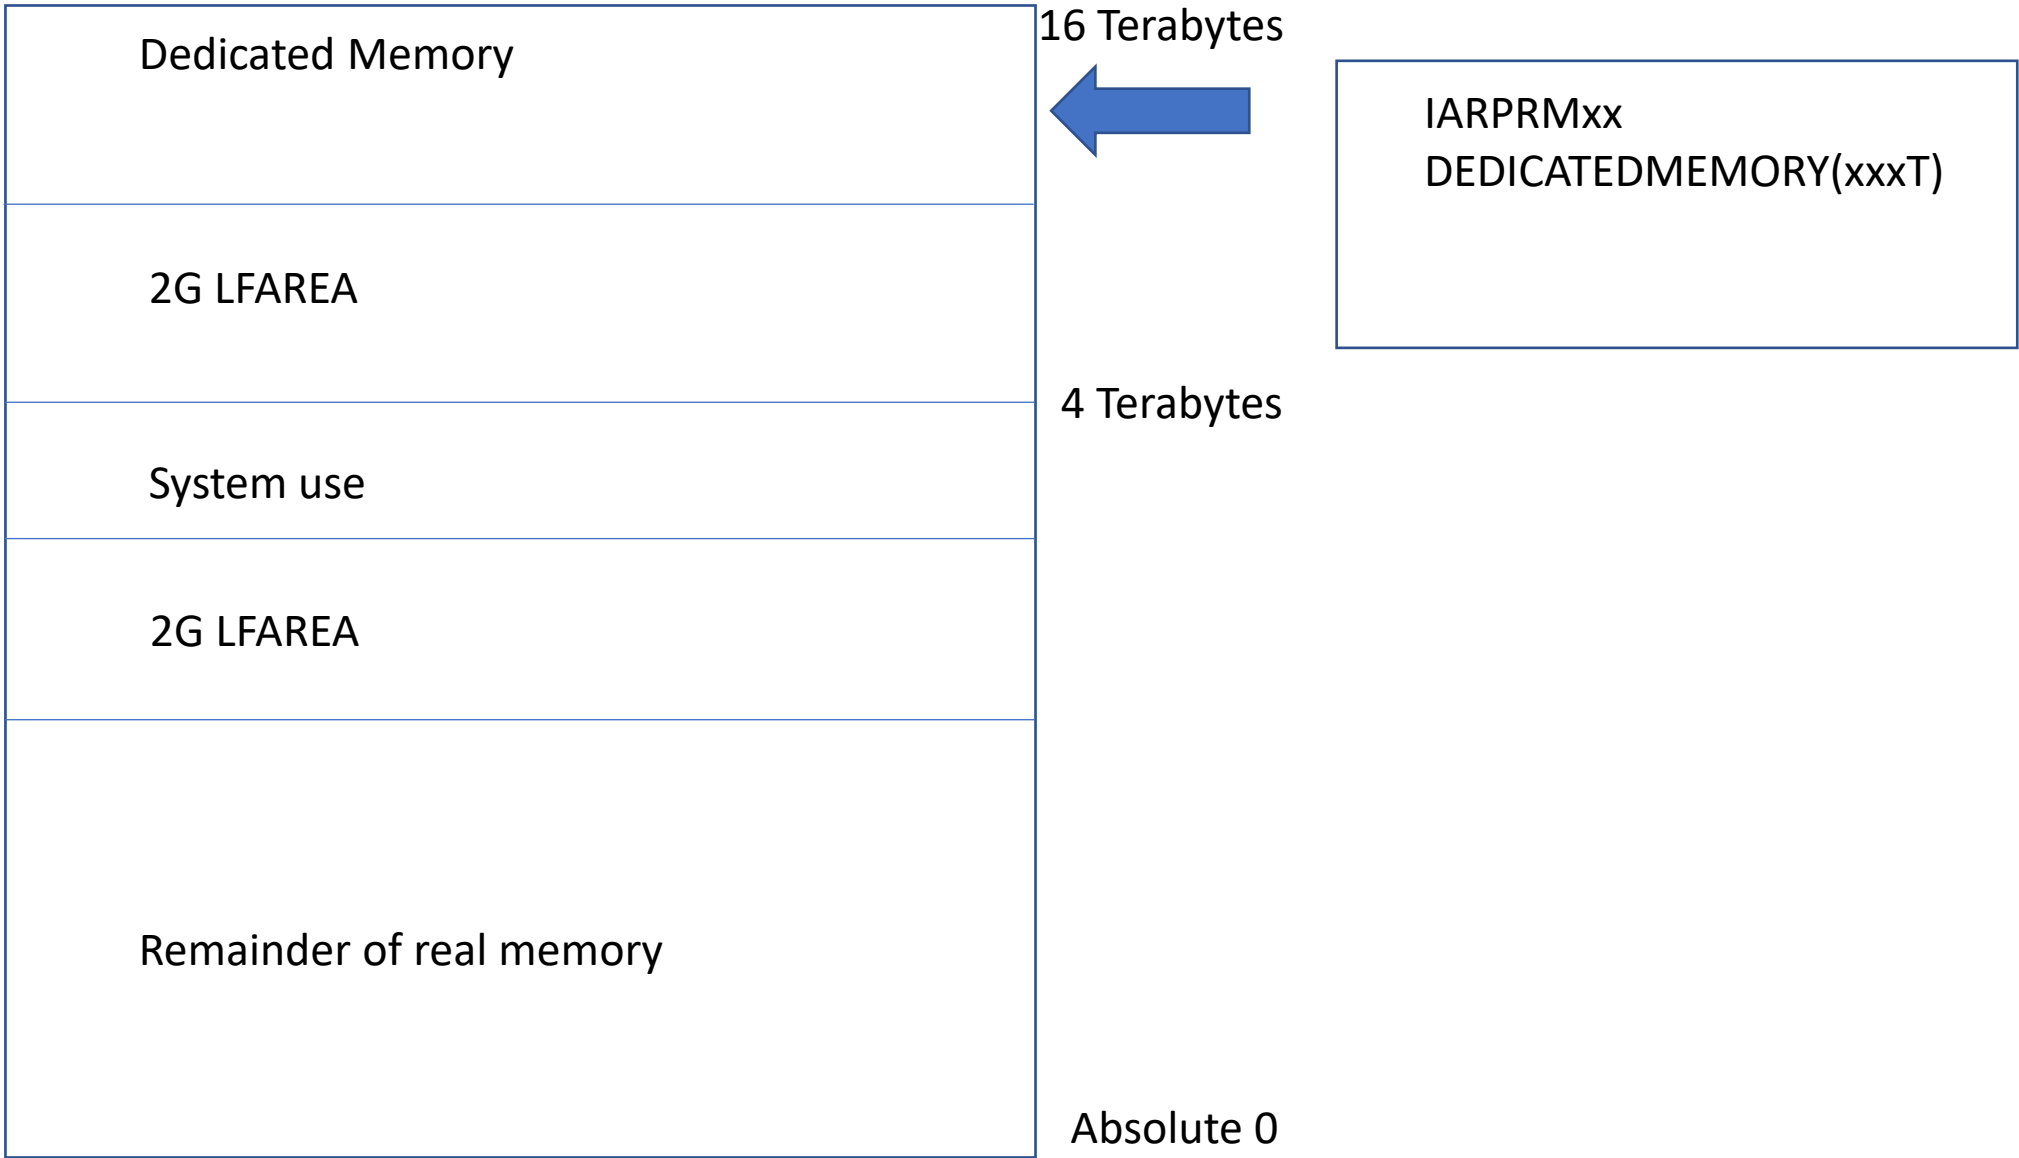
Initialization & IPL Considerations

- Historical real memory usage
 - SMF30HVR –High water mark of the number of real storage frames that are used to back 64-bit private storage.
 - SMF30HVA –High water mark of the amount of auxiliary storage that is used to back 64-bit private storage.
 - SMF30_InUseAs2GHWM –High water mark of the number of 2G frames in use by the job step (new in R3.1)
 - SMF30_NUM2GFailed -Number of 2G frames that could not be obtained because none were available at the time of the IARV64 GETSTOR request (new in R3.1)
- Possible estimate (over-estimate)
 - $((\text{SMF30HVR} + \text{SMF30HVA}) * 4K + \text{SMF30_InUseAs2GHWM} * 2G)$

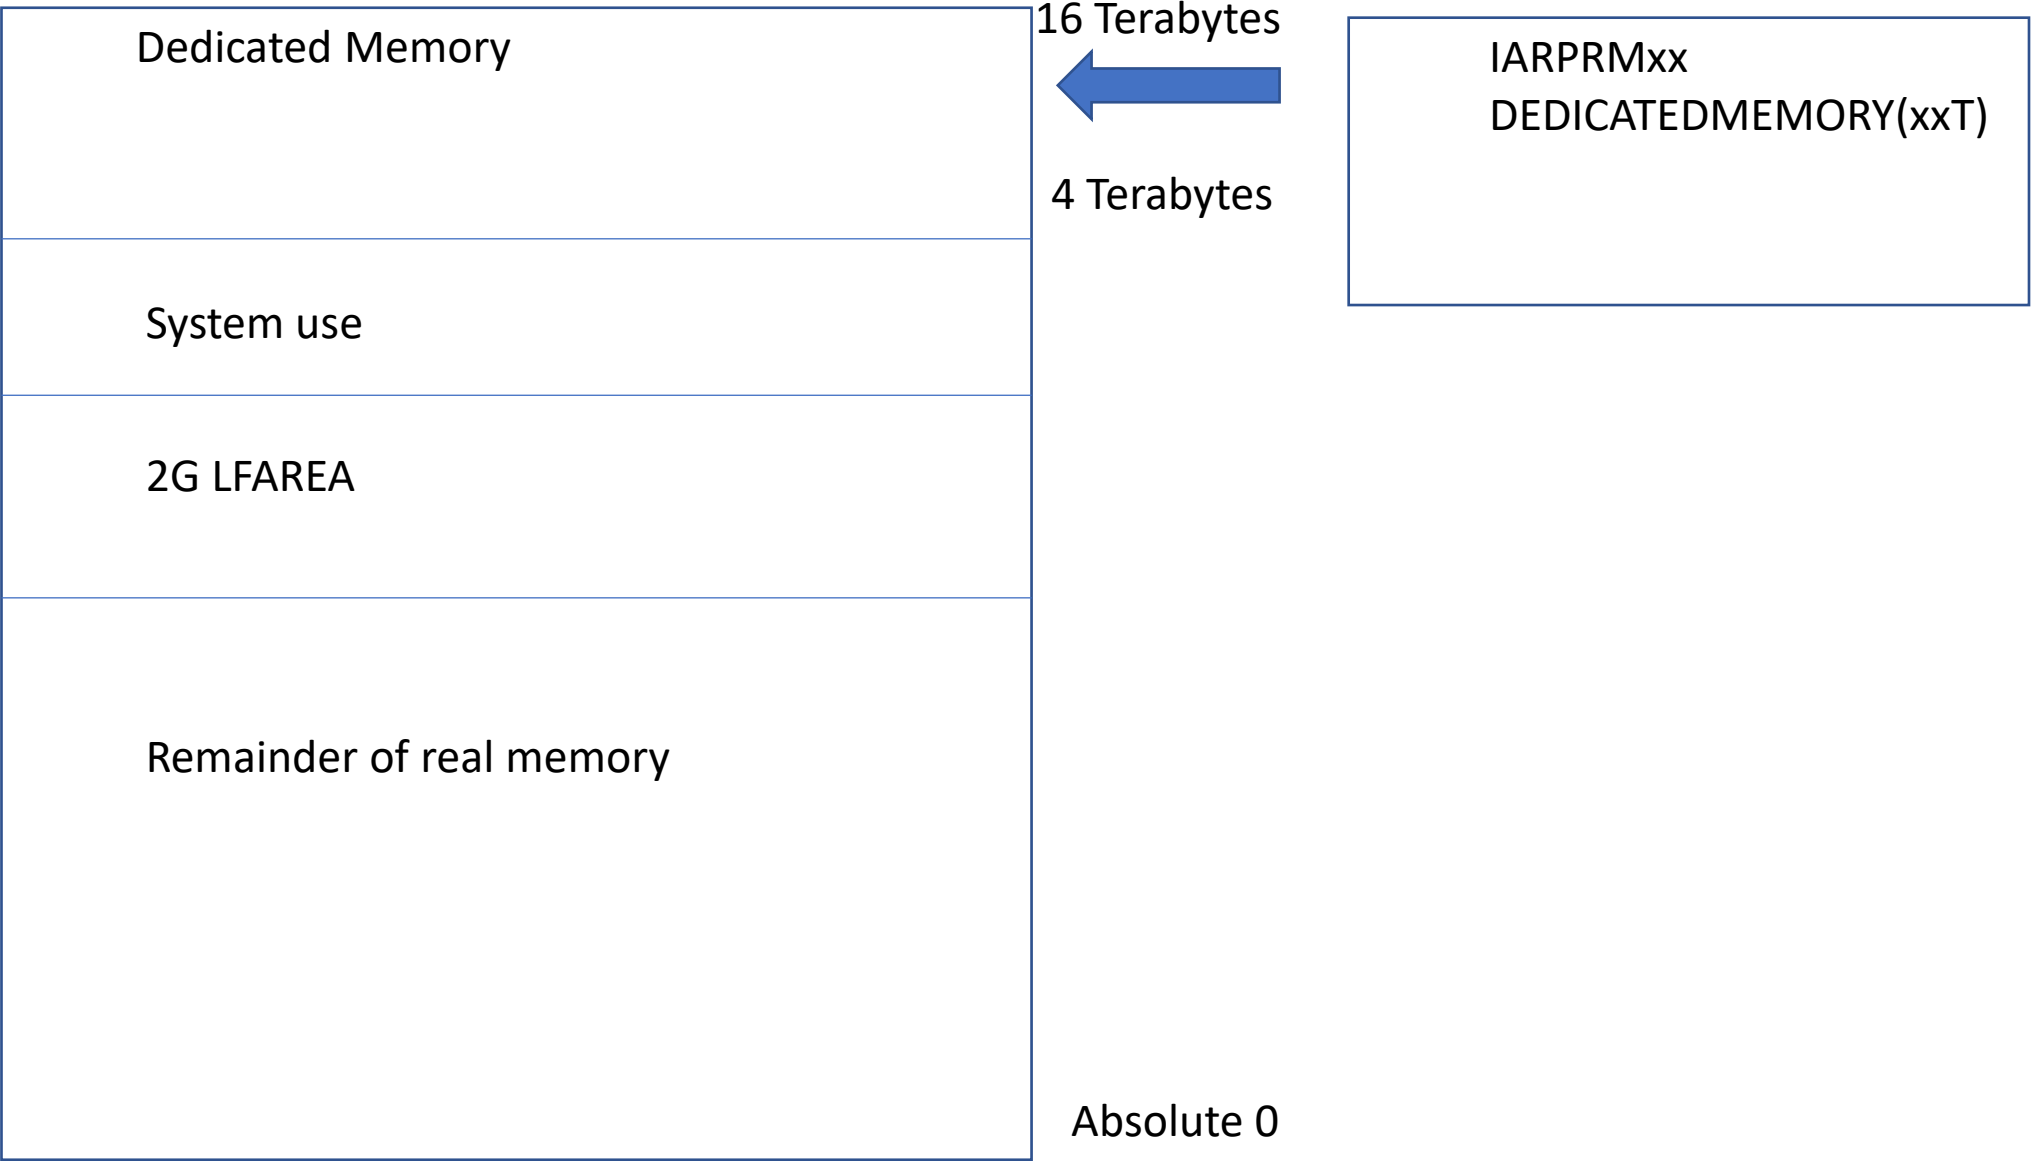
Initialization & IPL Considerations

- Some Dedicated Memory is used to manage Dedicated Memory and cannot be assigned to any job step
 - Referred to as system-assigned
- Approximately 2G out of every 124G of defined Dedicated Memory is system-assigned
- Example: If you want to define 400G of assignable Dedicated Memory then
 - $400/124 = 4$ (rounded up) – 408G should be requested.

Dedicated Memory and 2G LFAREA above 4T



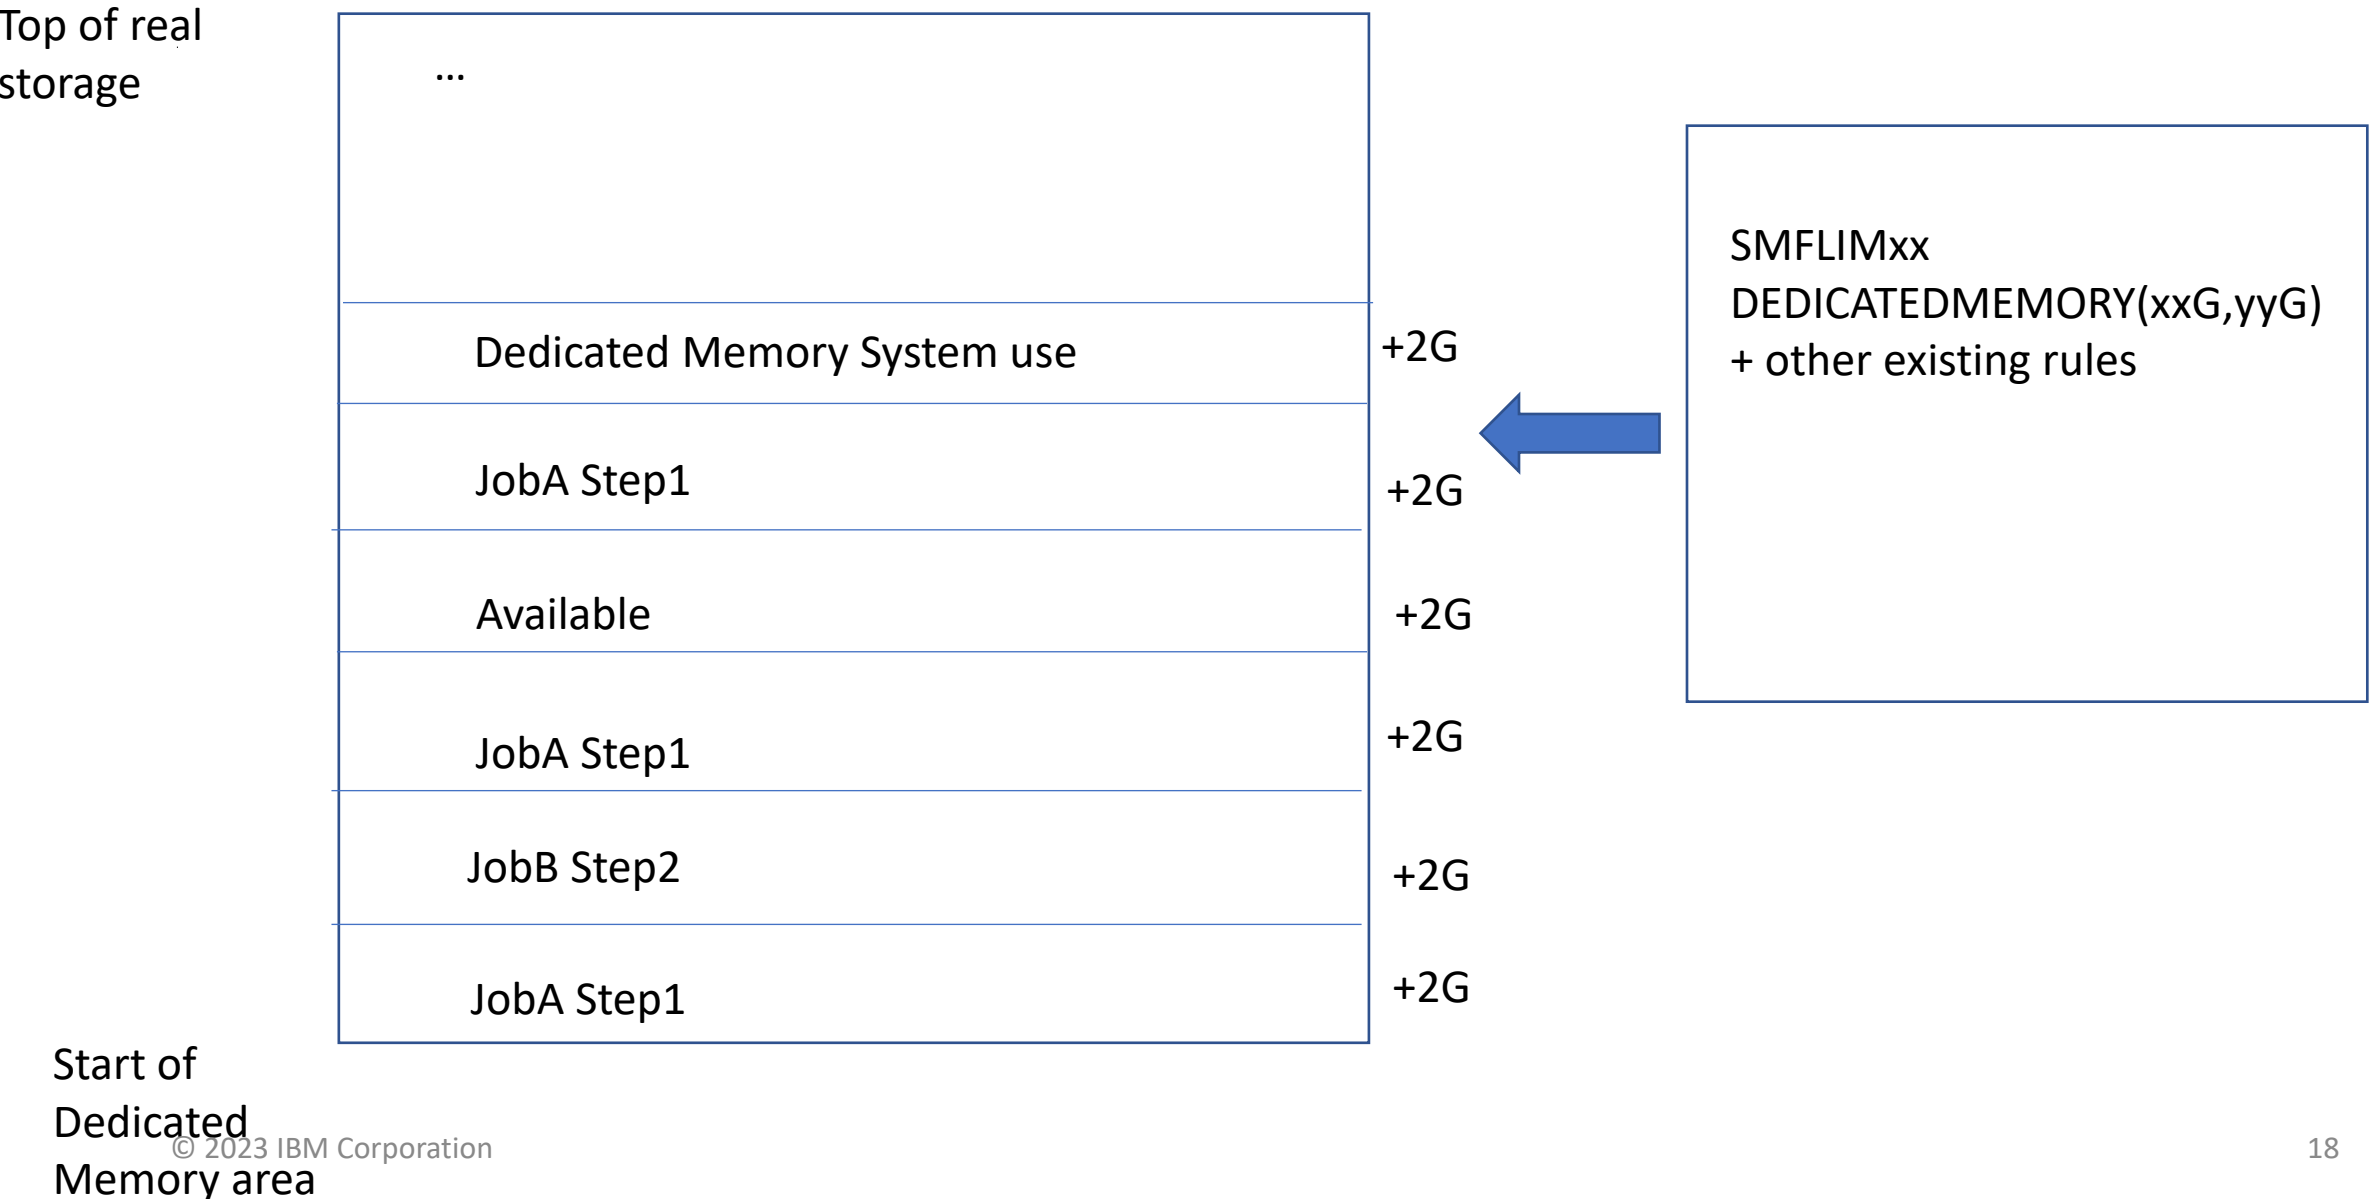
Just Dedicated Memory above 4T



Installation and Configuration

- Dedicated Memory is designated at the top of real storage
- A portion of Dedicated Memory is used by the system to manage dedicated memory.
- Dedicated memory is not completely ready for use until all of real storage is initialized as indicated by
 - message IAR060I (new) displayed
 - the new IAXVWAIT service
 - Allows an application to wait for real storage initialization to complete
 - RCERealFramesInitialized set to 1b
- System will “circle back” to those applications that requested Dedicated Memory when memory was still initializing
 - Once memory initialization is complete, assignments that were not completely satisfied during MSI will be satisfied.
- IAR060I REAL STORAGE INITIALIZATION COMPLETE. DEDICATED MEMORY AREA: xxG ONLINE, xxG OFFLINE

Dedicated Memory assignments



Initialization & IPL Considerations

- In order to assign Dedicated Memory to applications, parmlib member (SMFLIMxx) needs to be specified with the new DEDICATEDMEMORY parameter:
 - DEDICATEDMEMORY(0G,8G)
 - First value is the minimum amount of dedicated memory to assign
 - Second value is the target amount of dedicated memory to assign
 - If the minimum could not be satisfied, the application will be cancelled unless the application is a system address space
 - See the addendum for syntax rules and example messages.
- The amount of assigned dedicated memory to a job step can be dynamically changed after IPL through the SET SMFLIM command.
 - Note that if the application had started before the SET SMFLIM command, the change in size will not take effect until the application is restarted.

Initialization & IPL Considerations

Example of assigning up to 4G of Dedicated Memory to DUMPSRV:

REGION

JOBNAME(DUMPSRV)

DEDICATEDMEMORY(0G,4G)

Initialization & IPL Considerations

```
REGION SYSNAME(AQTS) JOBNAME(BIGSORT*) STEPNAME(GREEDY)  
DEDICATEDMEMORY(100G) MEMLIMIT(100G)
```

```
REGION SYSNAME(AQTS) JOBNAME(BIGSORT*) STEPNAME(LEAN)  
DEDICATEDMEMORY(0G,10G) MEMLIMIT(10G)
```

Expected result: Say that step GREEDY goes first. If the system has 100G of dedicated memory it is provided to GREEDY, otherwise the job is cancelled. GREEDY runs to completion. LEAN job step only requires 10G of dedicated memory. Job step initialization frees 90G before LEAN runs.

Initialization & IPL Considerations

- zCX initialization has changed to take Dedicated Memory into account when deciding whether the container can start
 - Previously just uses results from SYSEVENT STGTEST to determine the impact to the system
 - SYSEVENT STGTEST just reports on non-Dedicated Memory
- SVC dump processing has changed to capture to high virtual storage in the OPTIMIZE=NO path
 - Already captures to high virtual in the OPTIMIZE=YES path
 - Ensures that Dedicated Memory is used when Dedicated Memory is assigned to DUMPSRV

Installation and Configuration

```
IAR013I 4G STORAGE IS RECONFIGURABLE
IAR073I MEMORY CONFIGURATION
TOTAL MEMORY: 64G
ONLINE MEMORY: 64G
----- REQUESTED AMOUNTS -----
DEDICATEDMEMORY: 32G--IARPRMH2
LFAREA: *NOT REQUESTED*
RSU: 10%
----- ACTUAL AMOUNTS -----
DEDICATED MEMORY: 32G
ONLINE DEDICATED MEMORY: 32G
ASSIGNABLE DEDICATED MEMORY: 30G
2G LFAREA: 0M
2G LFAREA ABOVE 4T: 0M
1M LFAREA LIMIT: 0M
RECONFIGURABLE (RSU): 4G
ONLINE RECONFIGURABLE (RSU): 4G
IAR077A REPLY C TO CONTINUE WITH THE MEMORY CONFIGURATION OR CHANGE IT AND
RE-IPL.
```

Installation and Configuration

- The Dedicated Memory area is reconfigurable when the storage increment size is at least 2G
- Reconfiguration rules are different than for NON-Dedicated Memory
 - Only Dedicated Memory that is not assigned to a job and not in use by the system is eligible to be configured offline
 - All of Dedicated Memory can be configured offline if no job steps are assigned Dedicated Memory
- NOTE – As in 2.5, all memory reconfiguration is disabled if there is any non-Dedicated Memory above 4T (this has not changed)
 - RSU parameter is accepted if (1) there is no memory defined above 4T or (2) All the memory defined above 4T is Dedicated Memory
 - NOTE: RSU only applies to non-Dedicated Memory

Memory Reconfiguration

- Recall - three different types of memory reconfiguration:
 - By Range: CONFIG STOR(xx-yy),ONLINE/OFFLINE
 - By Amount: CONFIG STOR(xx),ONLINE/OFFLINE
 - By Element: CONFIG STOR(E=1),ONLINE/OFFLINE
- Changes for Dedicated Memory:
 - By Range – Ranges cannot span Dedicated Memory and NON-Dedicated Memory
 - By Amount: New DMEM keyword must be specified
 - CONFIG STOR(xx),DMEM,ONLINE/OFFLINE
 - A storage element with a mixture of Dedicated and NON-Dedicated Memory can be configured online or offline

D M=HIGH,DMEM

IEE174I changed to report on the range of Dedicated Memory

```
IEE174I 14.19.58 DISPLAY M          FRAME LAST  F      E  SYS=S7E0
DEDICATED MEMORY SIZE STATUS
STARTING ADDRESS IS 32G
ENDING ADDRESS IS 64G

IEE612I CN=RSMCON1  DEVNUM=03E0 SYS=S7E0
_

IEE163I MODE= RD
```

D M=STOR

- IEE174I changed to report on online ranges of Dedicated Memory
- D M=STOR,DMEM (new)
 - Reports which DMEM ranges are reconfigurable

D M=STOR

```

IEE174I 16.27.23 DISPLAY M                                FRAME LAST
REAL STORAGE STATUS
ONLINE-NOT RECONFIGURABLE
    0G-32G
ONLINE-RECONFIGURABLE
    32G-48G
ONLINE-DEDICATED MEMORY
    48G-60G
PENDING OFFLINE
    NONE
    0M IN OFFLINE STORAGE ELEMENT(S)
    4G UNASSIGNED STORAGE
STORAGE INCREMENT SIZE IS 4G

IEE612I CN=RSMCON1    DEVNUM=03E0  SYS=S7E0

IEE163I MODE= R

```

D M=STOR,DMEM

```
00- d m=stor,dmem
```

```
IEE174I 10.37.41 DISPLAY M          FRAME 1      F      E      SYS=S7E0
```

```
DEDICATED MEMORY STATUS
```

```
ONLINE-DEDICATED MEMORY, SOME SYSTEM ASSIGNED - NOT RECONFIGURABLE
```

```
60G-64G
```

```
ONLINE-DEDICATED MEMORY, SOME ASSIGNED TO A JOB - NOT RECONFIGURABLE
```

```
52G-64G
```

```
ONLINE-DEDICATED MEMORY - RECONFIGURABLE
```

```
32G-52G
```

```
PENDING OFFLINE
```

```
NONE
```

```
IEE612I CN=RSMCON1  DEVNUM=03E0 SYS=S7E0
```

```
-
```

```
IEE163I MODE= RD
```

Display Matrix with DMEM option

- ONLINE- DEDICATED MEMORY, SOME ASSIGNED TO SYSTEM – NOT RECONFIGURABLE
 - Some part of the increment is used to manage Dedicated Memory. Can be reconfigured offline only if the portion that it is used to manage goes offline.
- ONLINE-DEDICATED MEMORY, SOME ASSIGNED TO A JOB – NOT RECONFIGURABLE
 - Some part of the increment is assigned to a job
 - May not actually be used by the job
 - Cannot be configured offline until the job completes or the job step completes and the next jobstep does not require Dedicated Memory (or at least frees the memory associated with the increment being configured)
- ONLINE-DEDICATED MEMORY – RECONFIGURABLE
 - Ranges of Dedicated Memory that can be configured offline without cancelling a job.

D M=STOR,DMEM

```
CS/102 LIBRARY LOOKSIDE CRITICAL FAILURE (ABEND 3025 0000)  
* REASON=00000401)  
00- $HASP395 BPXAS      ENDED - RC=0000  
  - $HASP250 BPXAS PURGED -- (JOB KEY WAS DC322641)  
IEE174I 15.32.28 DISPLAY M          FRAME LAST    F      E    SYS=S7E0  
DEDICATED MEMORY STATUS  
ONLINE-DEDICATED MEMORY - RECONFIGURATION DISABLED  
    32G-64G  
    0M IN OFFLINE STORAGE ELEMENT(S)  
    0M UNASSIGNED STORAGE  
STORAGE INCREMENT SIZE IS 1G  
  
      -  
  
IEE612I CN=RSMCON1  DEVNUM=03E0 SYS=S7E0  
  
IEE163I MODE= RD
```

D M=STOR,DMEM

```
D M=STOR,DMEM
IEE174I 17.35.19 DISPLAY M 250
DEDICATED MEMORY STATUS
ONLINE-DEDICATED MEMORY, SOME SYSTEM ASSIGNED - NOT RECONFIGURABLE
    NONE
ONLINE-DEDICATED MEMORY, SOME ASSIGNED TO A JOB - NOT RECONFIGURABLE
    NONE
ONLINE-DEDICATED MEMORY - RECONFIGURABLE
    NONE
PENDING OFFLINE
    48G-64G
PENDING STATUS DUE TO:
RANGE          ASID      JOBNAME
 56G-62G        29        RSJCEX06
 50G-56G        2A        RSJCEY06
 0M IN OFFLINE STORAGE ELEMENT(S)
 0M UNASSIGNED STORAGE
STORAGE INCREMENT SIZE IS 4G
```


D M=STOR(E)

```
- d m=stor(e)
```

```
IEE174I 10.50.20 DISPLAY M
```

```
FRAME LAST
```

```
F
```

```
E
```

```
SYS=S7E0
```

```
ADDITIONAL ELEMENT STATUS
```

```
0: DMEM=0M
```

```
1: DMEM=16G
```

```
IEE612I CN=RSMCON1 DEVNUM=03E0 SYS=S7E0
```

```
—
```

```
IEE163I MODE= RD
```

IAXDMEM

- SYSREXX Exec for querying memory information.
 - MODIFY AXR,IAXDMEM <parameters>
- New Parameters
 - DMEM
 - DMEM,JOBLIST
 - DMEM,JOBNAME=jjjjjjjj
 - DMEM,ASID=nnnn

F AXR,IAXDMEM DMEM

```
* END
00- f axr,iaxdmem dmem
    IAR067I DEDICATED MEMORY V1.0
        36.0GB : TOTAL SIZE
        0.0GB : OFFLINE SIZE
        26.0GB : UNASSIGNED
        2.0GB : SYSTEM USE
IEE612I CN=RSMCON1  DEVNUM=03E0  SYS=S7E0
-
IEE163I MODE= RD
```

F AXR,IAXDMEM DMEM,JOBLIST

```
      2.0GB : SYSTEM USE
00- @iaxdmem dmem,joblist
      IAR068I DEDICATED MEMORY V1.0
      JOBNAME  ASID      ASSIGNED      IN USE
      AS01      0020      8.0GB        8.0GB
IEE612I CN=RSMCON1  DEVNUM=03E0  SYS=57E0
—
IEE163I  MODE=  RD
```

F AXR,IAXDMEM D MEM,JOBNAME=xx

```
JOBNAME=AS01      ASID=0020
  8.0GB : ASSIGNED
  8.0GB : IN USE
  8.0GB : MAX IN USE
PAGEABLE 4K STATISTICS
  0.0MB : IN USE FOR PAGEABLE 4K PAGES
  0.0MB : MAX IN USE FOR PAGEABLE 4K PAGES
PAGEABLE 1M STATISTICS
  0.0MB : IN USE FOR PAGEABLE 1M PAGES
  0.0MB : MAX IN USE FOR PAGEABLE 1M PAGES
FIXED 1M STATISTICS
  0.0MB : IN USE FOR FIXED 1M PAGES
  0.0MB : MAX IN USE FOR FIXED 1M PAGES
FIXED 2G STATISTICS
  8.0GB : IN USE FOR FIXED 2G PAGES
  8.0GB : MAX IN USE FOR FIXED 2G PAGES
DAT TABLE STATISTICS
  0.0MB : IN USE FOR DAT TABLES
```

New APIs

- IARQUERY
 - Authorized service to obtain information about real storage usage
- IAXVWAIT
 - Unauthorized service to wait for memory initialization to complete
- IARXNDM
 - Authorized service to obtain system information about Dedicated Memory usage

Summary

- Dedicated Memory is an installation defined area of memory that can be assigned in 2G units at the job step level
- Can be defined above 4T
- Only used for High Virtual Private
- Avoids first come first serve – provides a way to ensure that applications have sufficient memory

Appendix

- z/OS MVS: Initialization and Tuning Guide
- z/OS MVS: Initialization and Tuning Reference
- z/OS MVS: System Commands
- z/OS MVS: Authorized Assembler Services Reference
- z/OS MVS: Assembler Services Reference
- z/OS MVS: Authorized Assembler Services Guide
- z/OS MVS: System Messages
- z/OS MVS: System Management Facilities

Addendum

- Additional material on SMFLIMxx syntax and messages
- Additional material on IARPRMxx syntax
- Additional material on IARQUERY, IAXVWAIT, and IARXNDM

SMFLIMxx syntax

- New REGION attribute DEDICATEDMEMORY(xxxxxu[,yyyyyu]) identifies how much dedicated memory to assign.
 - The first value (xxxxxu) is the minimum amount of dedicated memory to assign.
 - It can be 1-5 digits, followed by the number of units (G=gigabytes, T=terabytes, P=petabytes).
 - Any value from 0G to 16384P is valid, but it must be a multiple of 2G.
 - The second value (yyyyyu) is the target amount of dedicated memory to assign.
 - Like the first value, must be 1-5 digits, G/T/P units, any value from 0G-16384P, must be a multiple of 2G.
 - This value is optional, and if not specified, will be the same as the first value (the minimum amount.)
- The DEDICATEDMEMORY keyword can be combined with any of the existing filters (like JOBNAME) and attributes (like MEMLIMIT.)
 - If you share the SMFLIMxx parmlib member with pre-z/OS 3.1 systems, those systems will not recognize the DEDICATEDMEMORY keyword and will ignore any rule with that specified. In this situation, create separate rules for DEDICATEDMEMORY, or use the SYSNAME filter to only apply the rules specifying DEDICATEDMEMORY to the z/OS 3.1 systems.

SMFLIMxx messages

- Message IEF043I will indicate the steps that SMFLIM processing has performed, if it is not suppressed using the JOBMSG(SUPPRESS) attribute.

- Example of message when the job applied dedicated memory:

```
IEF043I Actions taken by SMFLIMxx parmlib policy for MYJOB MYSTEP  
Step Dedicated Memory changed to (00002G,00008G) by policy - SMFLIM00 0001
```

- Example of message when the job was cancelled due to insufficient dedicated memory:

```
IEF043I Actions taken by SMFLIMxx parmlib policy for MYJOB MYSTEP2  
Step cancelled due to insufficient Dedicated Memory value (00002G,00002G) by policy - SMFLIM00 0002
```

Other RSM messages

- IAR062I JOBNAME <jjjjjjjj> ASSIGNED ADDITIONAL DEDICATED MEMORY.TOTAL REQUESTED=<rrrrr>G|T|P, TOTAL ASSIGNED=<aaaaa>G|T|P
- IAR063I NO DEDICATED MEMORY WAS ASSIGNED. DEDICATED MEMORY REQUESTED=<nnnnnn>G|T, AVAILABLE=<nnnnnn>G|T
 - Job is cancelled (unless system address space)
- IAR064I xxxxG|T DEDICATED MEMORY ASSIGNED
- IAR065I JOB IS NOT ELIGIBLE FOR DEDICATED MEMORY
 - OMVS is not eligible

IARPRMxx syntax

- DEDICATEDMEMORY(mmG | nnT)
- DMEM(mmG|nnT)

Specifies the amount of real memory to be designated as Dedicated Memory where the suffixes G represents Gigabytes and T represents Terabytes. The amount must be a multiple of 2G and the total amount of configured memory on the system must equal a multiple of 2G. The system rounds the specified value up or down to a storage increment multiple. Dedicated Memory is designated from the top of Central Storage, regardless of whether the memory is online. At least 16G of online non-Dedicated Memory must exist after the Dedicated Memory designation, otherwise the request is rejected. Some amount of Dedicated Memory must be used by the system to manage Dedicated Memory; at least 4G must be requested.

IARPRMxx syntax

- PROMPT(NO|YES)

Specifies whether the system should prompt the operator to continue with the IPL after message IAR073I is issued. Message IAR073I describes the system's memory configuration, including the amount of Dedicated Memory.

SMF30/RAX fields

SMF Field name	Offset	Size/Format	RAX Field	Description
SMF30_DMEMRequested2G	256/ 100x	8/FIXED	RAX64_DMemRequested2G	Requested Dedicated Memory in 2G units at start of job step
SMF30_DMEMMinRequested2G	264/ 108x	8/FIXED	Rax64_DMemMinRequested2G	Requested minimum amount of Dedicated Memory in 2G units at start of Job Step.
SMF30_DMEMAssigned2G	272/ 110x	8/FIXED	RAX64_DmemAssigned2G	Assigned Dedicated Memory in 2G units at start of job step
SMF30_DMEMNumInUseAs2G	280/ 118x	8/FIXED	RAX64_DMEMIInUseAs2G	Number of 2G Dedicated Memory frames in use
SMF30_DMEMNumInUseAsFixed1M	288/ 120x	8/FIXED	Rax64_DMemInUseAs1MFixed	Number of Fixed 1M Dedicated Memory frames in use
SMF30_DMEMNumInuseAsPageable1M	296/ 128x	8/FIXED	Rax64_DMemInUseAs1MPageable	Number of pageable 1M Dedicated Memory frames in use

SMF30/RAX fields

SMF Field name	Offset	Size/Format	RAX Field	Description
SMF30_DMEMNumInUseAs4K	304/ 130x	8/FIXED	Rax64_DMemInUseAs4K	Number of 4K Dedicated Memory frames in use
SMF30_DMEMNumInUseAsDATTables	312/ 138x	8/FIXED	Rax64_DMemQuadGroups*4+Rax64_DMemPGTs	Number of 4K Dedicated Memory frames in use as DAT tables
SMF30_DMEMNumInUseAs4KHWM	320/ 140x	8/FIXED	Rax64_DMemInUseAs4KHWM	High water mark for the number of Dedicated real storage frames (4K units) that is used to back 64-bit private storage.
SMF30_DMEMNumInUseAsPageable1MHW WM	328/ 148x	8/FIXED	Rax64_DMemInUseAs1MPageableHWM	High water mark for the number of Dedicated real storage frames (1M units) that is used to back Pageable 1M pages
SMF30_DMEMNumInUseAsFIXED1MHW WM	336/ 150x	8/FIXED	Rax64_DMemInUseAs1MFixedHWM	High water mark for the number of Dedicated real storage frames (1M units) that is used to back Fixed 1M pages
SMF30_DMEMNumInUseAs2GHWM	344/ 158x	8/FIXED	RAX64_DMEMInUseAs2GHWM	High water mark for the number of Dedicated real storage frames (4K units) that is used to back 2G pages

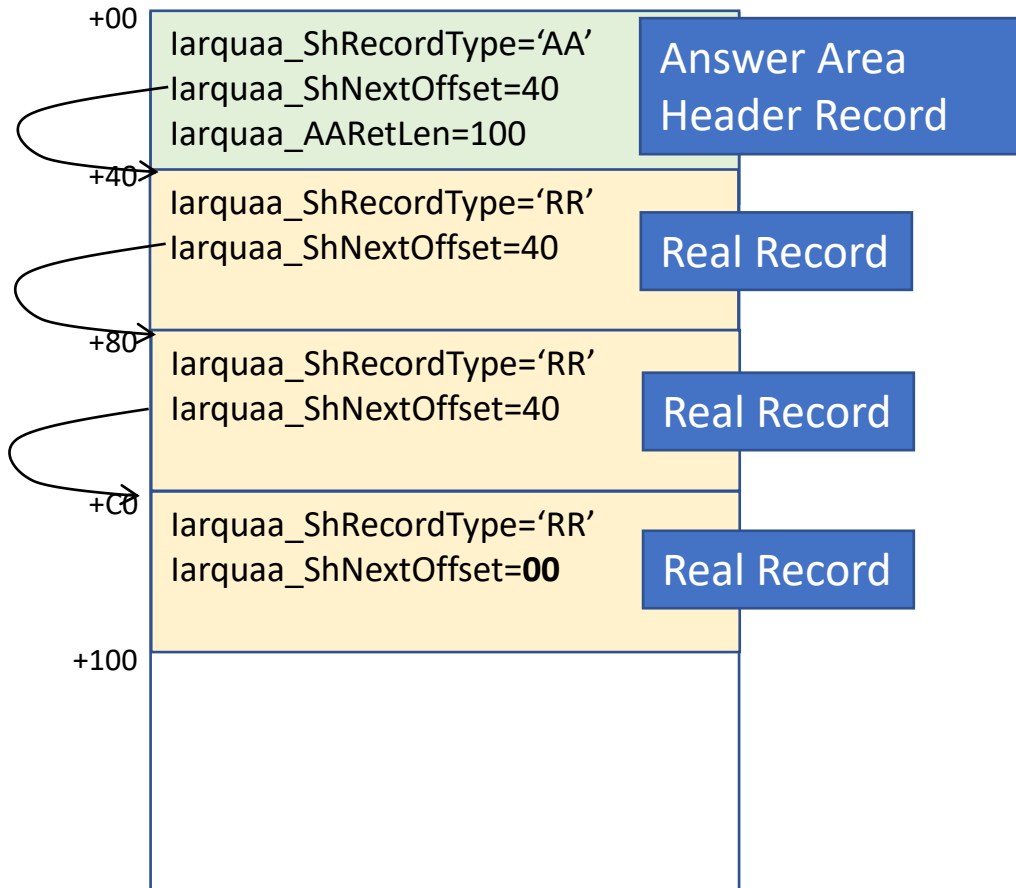
SMF30/RAX fields

SMF Field name	Offset	Size/Format	RAX Field	Description
SMF30_DMEMNumInUseAsDATTable sHWM	352/ 160x	8/FIXED	Rax64_DMENInUseAsDATTablesHWM	High water mark of the number of Dedicated real storage frames (4K units) that is used to back DAT tables
SMF30_DMEMNumInUseHWM	360/ 168x	8/FIXED	Rax64_DMemInUseHWM	High water mark of the number of Dedicated Memory frames in use in 4K units
SMF30_DMEMNum2GFailed	368/ 170x	8/FIXED	Rax64_DMem2GFailedCnt	Count of requested 2G frames eligible to backed by Dedicated Memory frames but were not due to availability
SMF30_DMEMNum1MFailed	376/ 178x	8/FIXED	Rax64_DMem1MFailedCnt	Count of requested 1M frames eligible to backed by Dedicated Memory frames but were not due to availability
SMF30_DMEMNum4KFailed	384 180x	8/FIXED	Rax64_DMem4KFailedCnt	Count of requested 4K frames eligible to backed by Dedicated Memory frames but were not due to availability
SMF30_NumInUseAs2GHWM	392/ 188x	8/FIXED	RAX64_2GPagesBackedInRealHWM	Number of 2G frames in use by the job step high water mark
SMF30_Num2GFailed	400/ 190x	8/FIXED	RAX64_2GFailedCnt	Number of 2G frames that could not be obtained because none were available at the time of the IARV64 request.

IARQUERY Service

- New IARQUERY service to obtain real storage information.
 - **IARQUERY REQINFO=REAL**
 - Return information about specific frames of real storage
 - **INSTARTADDR=,INCOUNT4K=**
 - Start of real storage range and size of range in 4K units
 - **OUTANSAREA=,INANSLEN=**
 - Address and length of output answer area
 - **[STOKEN=, [IncludeDS=YES/NO]]**
 - Filter by address space or data space STOKEN, include data spaces owned by address space
 - **RETCODE=,RSNCODE=**
 - Return and reason codes
- Authorized callers only, must be enabled for interrupts
- Most processing is done unserialized, and runs enabled

Answer Area Format



- Mapped by IARQUAA
- Answer Area Header Record
 - Contains overall length, record count, etc.
 - Resume address if output answer area fills
- Real Frame Record
 - Contains attributes of a frame of storage
 - Large frames are compressed into a single record
- Records are chained together by offset
 - Offset from start of current record to start of next record
 - Offset=0 indicates end of chain
- RC=4 with resume address if ANSAREA fills up before range is complete.

IARQUAA – Standard Header Fields

Field Name	Length	Description
larquaa_ShRecordType	2	Record Type
larquaa_ShVersion	1	Version
larquaa_ShLen	2	Length of this record
larquaa_ShNextOffset	2	Offset to next record

- Each Record in the answer area begins with a standard mapping.
 - This allows the record to be examined (and possibly skipped) without using the mapping for a specific record type
- PL/X uses ISA to define these fields so the larquaa_Shxxxx field names should always be used in PL/X (and ZTT).
- Assembler mapping defines the fields separately due to CBGEN restrictions with ISA fields.

IARQUAA – Answer Area Header Fields

Field Name	Length	Description
larquaa_AaRecordType	2	Record Type (“AA”)
larquaa_AaVersion	1	Version
larquaa_AaLen	2	Length of this record
larquaa_AaNextOffset	2	Offset to next record
larquaa_Aald	8	Eyecatcher “IARQUAA”
larquaa_Aa#Rec	8	Number of records returned
larquaa_Aa#Rem	8	Number of records not returned due to insufficient space in answer area
larquaa_AaRetLen	8	Length of data returned
larquaa_AaMaxLen	8	Total space needed to contain all requested data
larquaa_AAResumeAddr	8	Resume address if additional calls needed

IARQUAA – Real Record Fields

Field Name	Length	Description
larquaa_RrRecordType	2	Record Type (“RR”)
larquaa_RrVersion	1	Version
larquaa_RrLen	2	Length of this record
larquaa_RrNextOffset	2	Offset to next record
larquaa_RrRSA	8	Real Storage Address
larquaa_RrVSA	8	Virtual Storage Address
larquaa_RrShareToken	8	Share token for shared storage
larquaa_RrRangeSize	4	Number of 4K units represented by this record
larquaa_RrOwningAsid	2	Owning ASID
larquaa_RrFixCount	4	Fix Count

IARQUAA – Real Record Fields (continued)

Field Name	Length	Description
larquaa_RrAged	bit	Frame is aged
larquaa_RrDref	bit	Frame is DREF
larquaa_RrDataspace	bit	Dataspace Frame
larquaa_RrIOInProgress	bit	Paging I/O in progress
larquaa_RrTransition	bit	Frame may be in a transitional state
larquaa_RrFrameType	2	Frame Type Flags
larquaa_RrFrameUsage	2	Frame Usage
larquaa_RrFrameSize	1	Frame size (4k, 1M, 2G, quadframe, etc.)
larquaa_RrKey	bit(4)	Storage key if available
larquaa_RrFP	bit	Fetch protect status if available
larquaa_RrKeyFpValid	bit	Storage key and fetch protect status were returned

IARQUAA – IARQUAA_RrFrameType Flags

Field Name	Length	Description
larquaa_FtOffline	bit	Frame is offline
larquaa_FtPref	bit	Frame is preferred storage
larquaa_FtVirtEqReal	bit	V=R frame
larquaa_FtReconfigurable	bit	Frame is reconfigurable
larquaa_FtPermanent	bit	Permanent storage
larquaa_FtDatArea	bit	Frame is in reserved DAT frame area (Quad area)
larquaa_Ft2GArea	bit	Frame is in 2G LFAREA
larquaa_FtBadFrame	bit	Bad Frame
larquaa_FtDmemArea	bit	Frame is in Dedicated Memory area

IARQUAA – IARQUAA_RrFrameUsage Types

Field Name	Field Name	Field Name
larquaa_kUsageNotApplicable	larquaa_kUsageAvailable	larquaa_kUsageSystem
larquaa_kUsageNucleus	larquaa_kUsageFixedLPA	larquaa_kUsageFlawed
larquaa_kUsageUninitialized	larquaa_kUsageInTransition	larquaa_kUsagePftSpace
larquaa_kUsagePftDatTable	larquaa_kUsage31SQA	larquaa_kUsageCommon
larquaa_kUsageSharedGroup	larquaa_kUsageSharedSegment	larquaa_kUsageHVShared
larquaa_kUsagePrivate	larquaa_kUsage31LSQA	larquaa_kUsageDataSpace
larquaa_kUsageVioCache	larquaa_kUsage2GPage	larquaa_kUsageFreemained
larquaa_kUsageDatTable	larquaa_kUsageDeferDelete	larquaa_kUsageAssigned

Frame usage is mainly a translation of the PFTQID and PFTPGTYP fields, although other PFT values may factor in as well. There is no direct 1-1 correspondence of the values.

IARQUERY Return Codes

- **IARQUAA_kRcOK (0)**
 - IARQUERY request completed successfully
- **IARQUAA_kRcWarn (4)**
 - IARQUERY request completed successfully, however warning has been issued.
 - RSN=**IARQUAA_kRsnAnsAreaFull** (xx0402xx) indicates that the request was partially satisfied but the provided answer area is full
- **IARQUAA_kRcParmError (8)**
 - IARQUERY request has parameters in error.
- **IARQUAA_kRcEnvError (12)**
 - IARQUERY request has an environment error.
 - For example, a disabled caller
- **IARQUAA_kRcCompError (16)**
 - IARQUERY request has encountered a component error.

IAXVWAIT Service

- IAXVWAIT – Service to wait for memory to initialize
 - Can be LINK'd to
 - P=H=S, Task mode, unauthorized, AMODE 31, Primary ASC, no FRRs set
 - Register1 = 0 on input

IARXNDM service

- IARXNDM – Service to obtain Dedicated Memory system usage statistics
 - Callable service
 - Caller must include IAXEPVT for the address of the service
 - Key 0 Supervisor state, unlocked
 - Takes parameter mapped by IAXXNDM in Register 1

XNDM input fields and return codes

Field Name	Length	Description
XNDM_Id	8	The id of the control block
XNDM_Version	1	Version of the control block
Reserved space	7	Reserved space – must be 0

Return code (reg 15)	Reason code (reg 0)	Meaning
0	N/A	Successful completion
8	xx0001xx	XNDM_Id incorrect
8	xx0002xx	Incorrect version
8	xx0003xx	Reserved fields not zero
16	N/A	System failure

XNDM output fields

Field Name	Length	Description
XNDM_NumberOfJobsUsingDMEM	8	Total number of jobs using Dedicated Memory
XNDM_Total2G	8	Total amount of Dedicated Memory defined in 2G units
XNDM_TotalOnline2G	8	Total amount of online Dedicated Memory in 2G units
XNDM_Assignable2G	8	Total amount of assignable Dedicated Memory in 2G units
XNDM_Available2G	8	Total available (unassigned) Dedicated Memory in 2G units
XNDM_Requested2G	8	Total amount of requested Dedicated Memory in 2G units
XNDM_MinRequested2G	8	Total minimum requested Dedicated Memory in 2G units
XNDM_Assigned2G	8	Total assigned (to a job step) Dedicated Memory in 2G units
XNDM_InUseAs2G	8	Total amount of Dedicated Memory used as 2G frames
XNDM_InUseAs1MFixed	8	Total amount of Dedicated Memory used as 1M fixed frames
XNDM_InUseAs1MPageable	8	Total amount of Dedicated Memory used as 1M pageable frames (IARV64 GETSTORE PAGEFRAMESIZE=PAGEABLE1MEG)
XNDM_InUseAs4K	8	Total amount of Dedicated Memory used as 4K frames
XNDM_InUseAsDATTables4K	8	Total amount of Dedicated Memory used to back DAT tables (system use)

z/OS 3.1 IBM Education Assistant

Solution Name: z/OS Data Gatherer support for Dedicated Memory

Solution Element(s): z/OS Data Gatherer



Overview – z/OS Data Gatherer support for dedicated memory

Dedicated memory usage data is collected by z/OS Data Gatherer and new metrics are implemented in:

- SMF record type 71, Paging Data Section (system scope)
- Monitor III measurement tables
 - ERBGEIG3 - General information table (system scope)
 - ERBASIG3 - Address space identification table (address space scope)

This support is included in the GA shipment of the z/OS 3.1 Data Gatherer (HRG77E0) deliverable.

Usage & Invocation: SMF record type 71 (1)

Monitor I Data Gatherer

- invokes RSM service IARXNDM to retrieve summarized dedicated memory usage data to obtain a system level view
- New metrics are stored in Paging Data Section of SMF71 records
- SMF71 record level (SMF71SRL) is increased to x'8F'

Offsets	Name	Length	Format	Description
1332 534	SMF71RFL	4	binary	Flags: Bit Meaning when set ... 2 Dedicated Memory usage counts are invalid 3-31 Reserved
...				
2440 988	SMF71_DMemAssignable2G	8	floating	Total amount of dedicated memory at system initialization that can be used by address spaces in 2G units. Note: This field does not include dedicated memory used by the system.
2448 990	SMF71_DMemNumberofJobsUsingDmem_M	8	floating	Minimum number of address spaces using dedicated memory.
2456 998	SMF71_DMemNumberofJobsUsingDmem_X	8	floating	Maximum number of address spaces using dedicated memory.
2464 9A0	SMF71_DMemNumberofJobsUsingDmem_A	8	floating	Average number of address spaces using dedicated memory.

Usage & Invocation: SMF record type 71 (2)

Offsets	Name	Length	Format	Description
2472 9A8	SMF71_DMemTotal2G_M	8	floating	Minimum amount of online and offline dedicated memory in 2G units, including dedicated memory used by the system.
2480 9B0	SMF71_DMemTotal2G_X	8	floating	Maximum amount of online and offline dedicated memory in 2G units, including dedicated memory used by the system.
2488 9B8	SMF71_DMemTotal2G_A	8	floating	Average amount of online and offline dedicated memory in 2G units, including dedicated memory used by the system.
2496 9C0	SMF71_DMemTotalOnline2G_M	8	floating	Minimum amount of online dedicated memory in 2G units, including dedicated memory used by the system.
2504 9C8	SMF71_DMemTotalOnline2G_X	8	floating	Maximum amount of online dedicated memory in 2G units, including dedicated memory used by the system.
2512 9D0	SMF71_DMemTotalOnline2G_A	8	floating	Average amount of online dedicated memory in 2G units, including dedicated memory used by the system.
2520 9D8	SMF71_DMemAvailable2G_M	8	floating	Minimum amount of available dedicated memory in 2G units that can be used by address spaces.
2528 9E0	SMF71_DMemAvailable2G_X	8	floating	Maximum amount of available dedicated memory in 2G units that can be used by address spaces.
2536 9E8	SMF71_DMemAvailable2G_A	8	floating	Average amount of available dedicated memory in 2G units that can be used by address spaces.
2544 9F0	SMF71_DMemRequested2G_M	8	floating	Minimum amount of requested dedicated memory in 2G units over all address spaces
2552 9F8	SMF71_DMemRequested2G_X	8	floating	Maximum amount of requested dedicated memory in 2G units over all address spaces
2560 A00	SMF71_DMemRequested2G_A	8	floating	Average amount of requested dedicated memory in 2G units over all address spaces

Usage & Invocation: SMF record type 71 (3)

Offsets	Name	Length	Format	Description
2568 A08	SMF71_DMemMinRequested2G_M	8	floating	Lowest minimum amount of requested dedicated memory in 2G units over all address spaces
2576 A10	SMF71_DMemMinRequested2G_X	8	floating	Highest minimum amount of requested dedicated memory in 2G units over all address spaces
2584 A18	SMF71_DMemMinRequested2G_A	8	floating	Average minimum amount of requested dedicated memory in 2G units over all address spaces
2592 A20	SMF71_DMemAssigned2G_M	8	floating	Minimum amount of assigned dedicated memory over all address spaces
2600 A28	SMF71_DMemAssigned2G_X	8	floating	Maximum amount of assigned dedicated memory over all address spaces
2608 A30	SMF71_DMemAssigned2G_A	8	floating	Average amount of assigned dedicated memory over all address spaces
2616 A38	SMF71_DMemInUseAs2G_M	8	floating	Minimum amount of dedicated memory used as 2G frames over all address spaces
2624 A40	SMF71_DMemInUseAs2G_X	8	floating	Maximum amount of dedicated memory used as 2G frames over all address spaces
2632 A48	SMF71_DMemInUseAs2G_A	8	floating	Average amount of dedicated memory used as 2G frames over all address spaces
2640 A50	SMF71_DMemInUseAs1MFixed_M	8	floating	Minimum amount of dedicated memory used as 1M fixed frames over all address spaces
2648 A58	SMF71_DMemInUseAs1MFixed_X	8	floating	Maximum amount of dedicated memory used as 1M fixed frames over all address spaces
2656 A60	SMF71_DMemInUseAs1MFixed_A	8	floating	Average amount of dedicated memory used as 1M fixed frames over all address spaces

Usage & Invocation: SMF record type 71 (4)

Offsets	Name	Length	Format	Description
2664 A68	SMF71_DMemInUseAs1MPageable_M	8	floating	Minimum amount of dedicated memory used as 1M pageable frames over all address spaces
2672 A70	SMF71_DMemInUseAs1MPageable_X	8	floating	Maximum amount of dedicated memory used as 1M pageable frames over all address spaces
2680 A78	SMF71_DMemInUseAs1MPageable_A	8	floating	Average amount of dedicated memory used as 1M pageable frames over all address spaces
2688 A80	SMF71_DMemInUseAs4K_M	8	floating	Minimum amount of dedicated memory used as 4K frames over all address spaces
2696 A88	SMF71_DMemInUseAs4K_X	8	floating	Maximum amount of dedicated memory used as 4K frames over all address spaces
2704 A90	SMF71_DMemInUseAs4K_A	8	floating	Average amount of dedicated memory used as 4K frames over all address spaces
2712 A98	SMF71_DMemInUseAsDATTables4K_M	8	floating	Minimum number of frames in 4K units backing DAT tables (system use)
2720 AA0	SMF71_DMemInUseAsDATTables4K_X	8	floating	Maximum number of frames in 4K units backing DAT tables (system use)
2728 AA8	SMF71_DMemInUseAsDATTables4K_A	8	floating	Average number of frames in 4K units backing DAT tables (system use)

Usage & Invocation: Monitor III measurement table ERBGEIG3 (1)

Monitor III Data Gatherer collects dedicated memory usage data

- having system scope from RCE / RCE64
 - New metrics are stored in General information table (GEIG3)
 - GEIG3 version number (GEIVERG3) is increased to x'1A'

Offsets	Name	Length	Format	Description
832 340	GEI_DMemAssignable2G	8	floating point	Total amount of dedicated memory at system initialization that can be used by address spaces in 2G units. Note: This field does not include dedicated memory used by the system.
840 348	GEI_DMemTotalOnline2G	8	floating point	Accumulated amount of online dedicated memory in 2G units, including dedicated memory used by the system. ¹
848 350	GEI_DMemTotal2G	8	floating point	Accumulated amount of online and offline dedicated memory in 2G units, including dedicated memory used by the system. ¹
856 358	GEI_DMemAFC2G	8	floating point	Accumulated amount of available dedicated memory in 2G units that can be used by address spaces. ¹
¹ Sum of values obtained at each sample. To obtain average values, divide by the number of valid samples (SSHSMPCR).				

Usage & Invocation: Monitor III measurement table ERBASIG3 (1)

Monitor III Data Gatherer collects dedicated memory usage data

- having address space scope from RAX64
 - Address space identification table (ASIG3) is enhanced to store new metrics
 - ASIG3 version number (ASIVERG3) is increased to x'1F'

Offsets	Name	Length	Format	Description
1248 4E0	ASI_DMemAFC	8	floating point	Amount of dedicated memory assigned to the address space in 4K units that is currently not in use ¹
1256 4E8	ASI_DMemAFC2G	8	floating point	Amount of dedicated memory assigned to the address space in 2G units that is currently not in use ¹
1264 4F0	ASI_DMemAFC1M	8	floating point	Amount of dedicated memory assigned to the address space in 1M units that is currently not in use ¹
1272 4F8	ASI_DMemAssigned2G	8	floating point	Amount of dedicated memory assigned to the address space in 2G units ¹
¹ Sum of values obtained at each sample. To obtain average values, divide by the number of valid samples (SSHSMPPNR).				

Usage & Invocation: Monitor III measurement table ERBASIG3 (2)

Offsets	Name	Length	Format	Description
1280 500	ASI_DMemInUseAs1MFixed	8	floating point	Amount of dedicated memory in use as 1M fixed frames ¹
1288 508	ASI_DMemInUseAs1MPageable	8	floating point	Amount of dedicated memory in use as 1M pageable frames ¹
1296 510	ASI_DMemInUseAs2G	8	floating point	Amount of dedicated memory in use as 2G fixed frames ¹
1304 518	ASI_DMemInUseAs4K	8	floating point	Amount of dedicated memory in use as 4K frames ¹
1312 520	ASI_DMemMinRequested2G	8	floating point	Minimum amount of dedicated memory requested by the address space in 2G units ¹
1320 528	ASI_DMemRequested2G	8	floating point	Amount of dedicated memory requested by the address space in 2G units ¹
1328 530	*	16	*	Reserved
¹ Sum of values obtained at each sample. To obtain average values, divide by the number of valid samples (SSHSMPCR).				

Appendix

- Documentation
 - z/OS Data Gatherer Programmer's Guide, GC31-5701
 - MVS System Management Facilities (SMF), SA38-0667